



# Image Classification

Barbaros Cetiner, PhD, PE

NHERI SimCenter

August 21, 2024

This material is based upon work supported by the National Science Foundation under Grant No. (1612843 & 2131111).

Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author(s) and do not necessarily reflect the views of the National Science Foundation.



# Outline

- Traditional Machine Learning Methods for Image Classification
- Deep Learning Methods
- Building Blocks of CNNs
- Attention Networks and Transformers
- Training a Deep Learning Model

# Traditional ML Methods for Image Classification




# Traditional Methods

Traditional image classification methods consist of two steps

1. Feature extraction (encoding)
2. Establishing a connection between features and image labels (decoding)


## Encoding



### Typical Algorithms:

- Scale invariant feature transform (SIFT)
- Histogram of oriented gradients (HOG)
- Binary robust invariant scalable keypoints (BRISK)

## Decoding



### Typical Algorithms:

- Support vector machines (SVM)
- Decision tress
- Neural networks

# Deep Learning Methods



# Traditional vs Deep Learning Methods

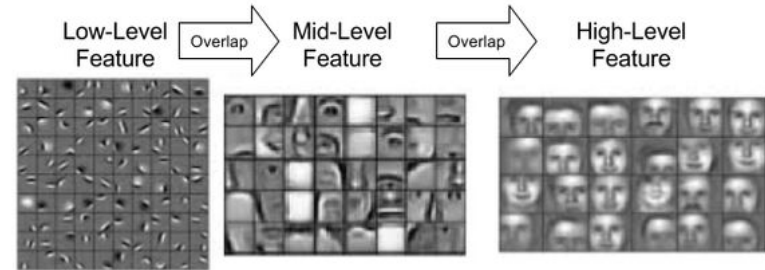
## Traditional Approaches

- Identify the features to solve a problem
- Develop methods to extract these features




## Deep Learning Methods

- Identifies the features it needs to solve a problem using the presented data



# Advantages of Deep Learning Methods

- End-to-end methods that require less feature engineering
- Consistently more accurate than traditional methods



Convolutional Neural Networks (CNN) are one of the most popular deep learning methods for image recognition.





# Building Blocks of CNNs





# Convolutional Layers

- Set of filters (or kernels), parameters of which are to be learned throughout the training
- Each filter convolves with the image and creates a feature (activation) map
- The kernel on the right has a **stride** of 1



# Pooling Layers

- Reduces the dimensionality of the computations
- Similar to convolutional layer, runs a filter across the entire input
- Unlike convolutional layer applies an aggregation function to the values
- What you see on the right is **Max pooling**




# Activation Functions

- **Note:** the composition of two linear functions is a linear function
- Activation functions enable nonlinear modeling


**Sigmoid**

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




**Leaky ReLU**

$$\max(0.1x, x)$$



**tanh**

$$\tanh(x)$$




**Maxout**

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

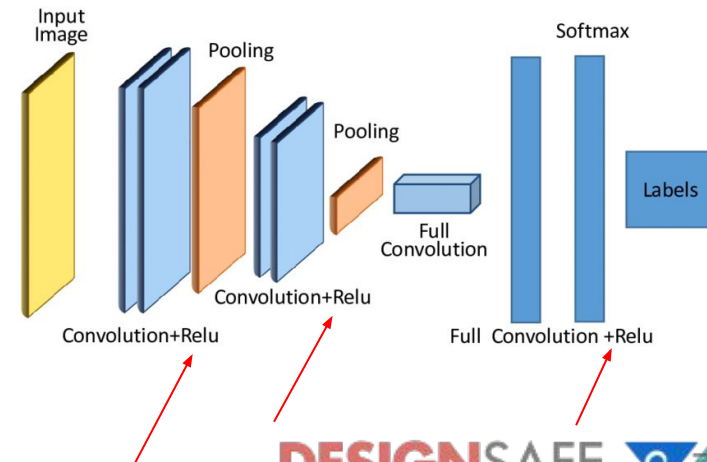
**ReLU**

$$\max(0, x)$$





**ELU**

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



# Dropout Layers

- Ensures that the model does not only rely on just the dominant features (avoids overfitting)
- Works by randomly killing some activations at each epoch
- **$p=0.5$**  kills half of the connections




# Quick Detour: What is Overfitting?


**Overfitting:** memorizing,  
instead of learning

Model remembers patterns,  
noise and random fluctuations,  
hence fail to perform well on  
unseen data (i.e., generalize)

**Tell-tale sign:** Divergence  
between training & validation  
accuracies (or losses)





Overfitting **✗**



Not overfitting **✓**

# Fully Connected (FC) Layers

- Also known as the Classifier Head or Dense Layer
- This layer performs classification based on the features extracted through the previous layers
- FC layers usually use SoftMax activation function to classify inputs




# Attention Networks and Transformers







# Why Use Attention-Based Models?

- CNNs do not capture the relative positions of different features
- The two figures below are the same to a CNN
- CNNs start at a single pixel and zoom out, attention-based models slowly bring the whole fuzzy image into focus

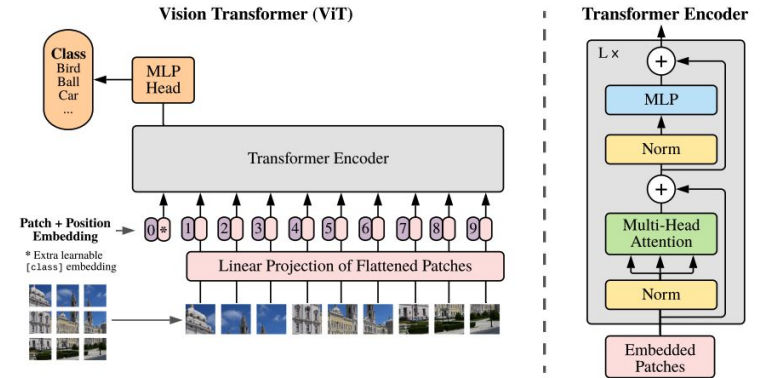


# Attention Networks




# Transformers


- Split an image into patches:
  - Image patches are treated the same way as tokens (words) in an NLP application.
- Flatten the patches
- Produce lower-dimensional linear embeddings from the flattened patches
- Add positional embeddings
- Feed the sequence as an input to a standard transformer encoder



# Training a Deep Learning Model for Image Recognition

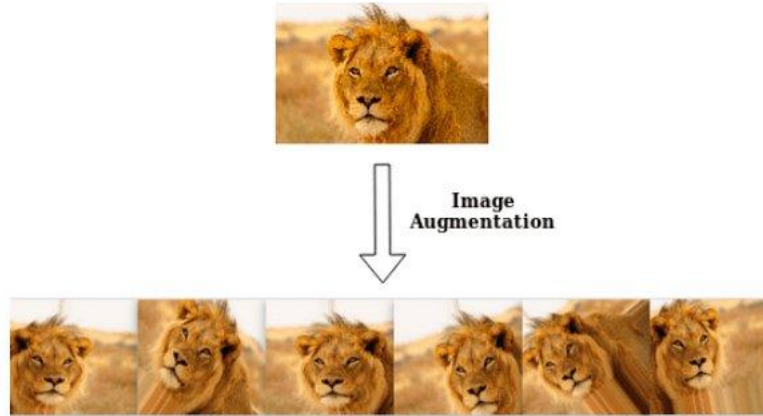


# Dataset Preparation: Class Imbalance



- **Always have a balanced training dataset** (i.e., equal number of samples for each class)
- Else, the training model will spend most of its time on the majority classes (and possibly have a bias towards these classes)

# Dataset Preparation: Image Augmentation



- Improves model performance by forming new and different examples to train datasets (expands the training dataset)
- Improves robustness of the model

# Dataset Preparation: Image Augmentation

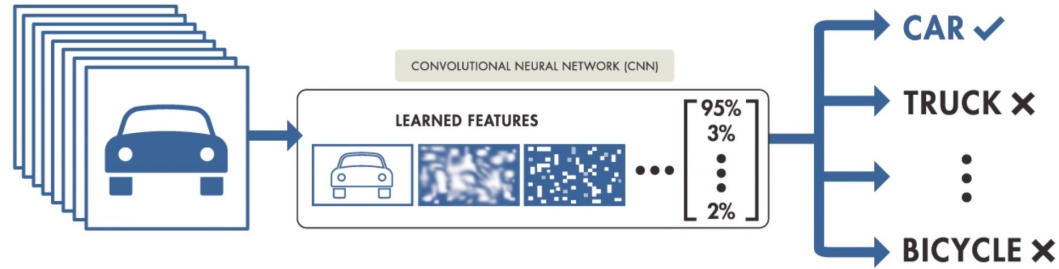


Popular techniques:

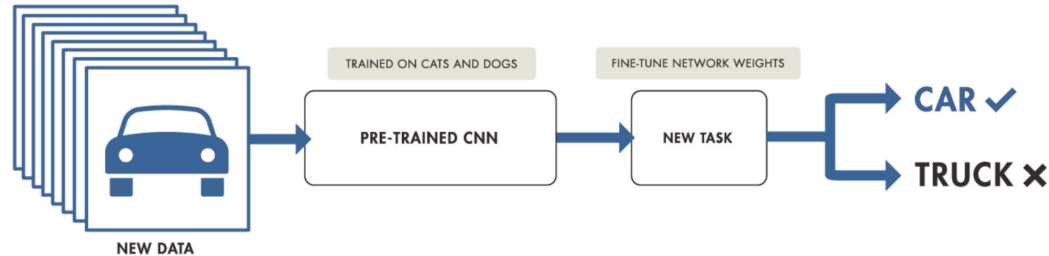
- Rotation,
- Gaussian noise,
- Crop,
- Hue and saturation adjustment,
- Elastic transform,
- Coarse dropout

# Model Training: Transfer Learning

## TRAINING FROM SCRATCH



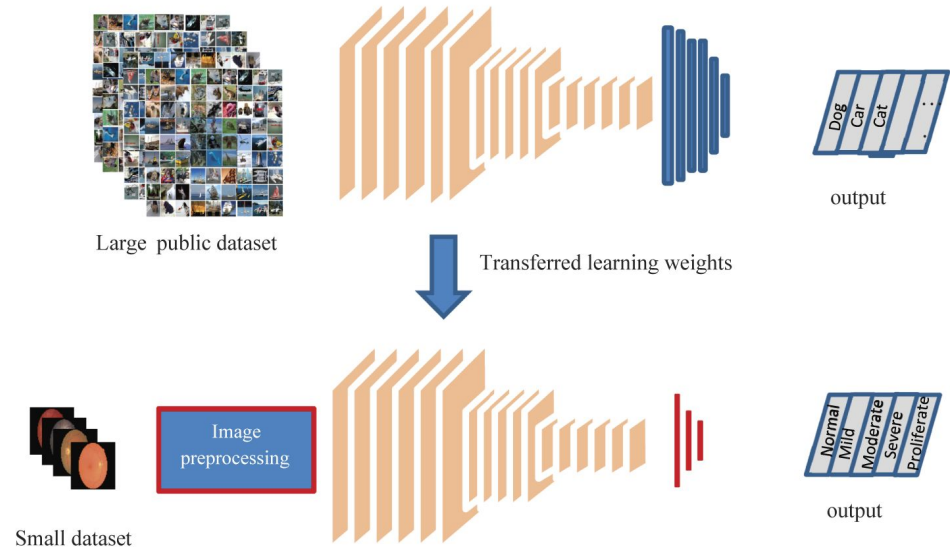
## TRANSFER LEARNING






# Model Training: Transfer Learning

- CNNs detect low level features in the first layer, forms in the middle layer, and task-specific features in the latter layers
- Transfer learning benefits from the lower features detected for a large dataset.
- Benefits:
  - Reduced training time,
  - Improved neural network performance in the absence of large training dataset



# Model Training: The Right Architecture

For civil engineering applications (small datasets) using an architecture with a good accuracy level on a standard dataset is the best!



# Model Training: Right Programming Approach

- Pytorch and Tensorflow are the most popular Python libraries for Deep Learning
- SimCenter's BRAILS provides easy-to-use end-to-end pipelines for training deep learning models
  - Little to no knowledge of deep learning is required to use BRAILS
  - Makes the difficult training decisions for the user



# Questions?

